

SENG 474 A02: Assignment 2

Model Selection Experiments: *Logistic Regression, Support
Vector Machines, and k-Fold Cross Validation*

Sean McAuliffe, V00913346

February 27th, 2023

Contents

1	Background	3
2	Environment and Implementation	3
3	k-Fold Cross Validation	3
4	Logistic Regression	5
5	Support Vector Machines	8
5.1	Finding the Optimal Regularization Parameter	9
6	Gaussian Kernel SVM	10

1 Background

This report summarizes the results of a series of experiments conducted which implement two methods for binary classification: Logistic Regression, and Support Vector Machines (SVM). The purpose of this exercise was to compare the performance of these two methods, as well as to practice the use of cross-entropy loss and k -Fold Cross Validation for selection of optimal model hyperparameters.

For each of these approaches a number of models were trained on the Fashion-MNIST [dataset](#). The Fashion-MNIST dataset is a collection of 70,000 grayscale 28x28 pixel images of clothing items in 10 classes. For the purposes of these experiments, only classes 0 and 6 (T-shirt and Shirt) were used. Furthermore, to reduce runtimes during SVM training, only 2,400 training images, and 2000 testing images were used throughout the experiments.

2 Environment and Implementation

Linear Regression and SVM implementation from **scikit-learn** were used for this assignment. Of note for this report, the **scikit-learn** library implements the *linear_model.LinearRegression* and *svm.SVC* classes for Logistic Regression and SVM, respectively. These classes were used to train and test the models for this assignment.

All code for this assignment, including data import, preprocessing model creation, training, evaluation and plot generation was done in Jupyter Notebooks running on a **Python 3.10.6** kernel. Relying on utility functions provided with the Fashion-MNIST dataset, the source code for which has been provided with included with the .ipynb files with this report.

3 k -Fold Cross Validation

The k -Fold Cross Validation method is an improvement on the simple cross-validation which was used to evaluate model performance in the previous assignment. Cross validation involves splitting the available data into

training and validation sets; models will be trained on the training set, and scored against the validation set. In this way the true risk of the model can be estimated on data which was not used to train the model.

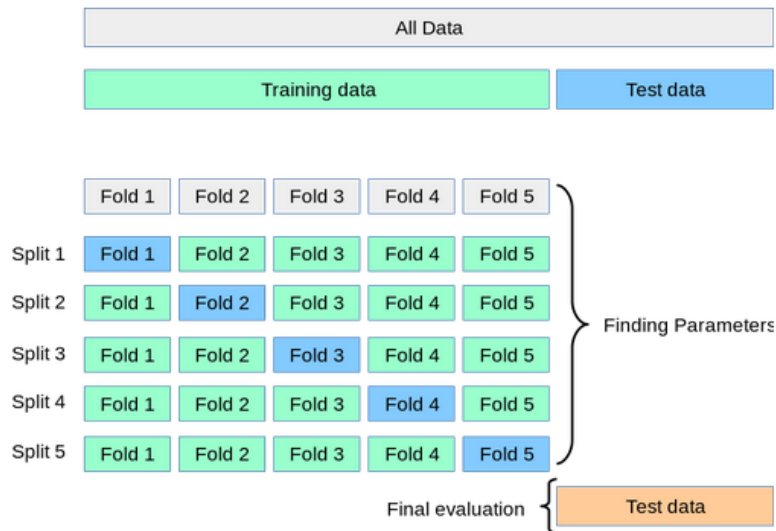


Figure 1: k-Fold Cross Validation Method, source: scikit-learn.org

The k -Fold Cross Validation method is a better way to evaluate model performance. It is more robust to the choice of training and validation sets, and is more computationally efficient than the simple method. The k -Fold Cross Validation method involves splitting the available data into k equal sized subsets. Then for each i between 0 and $k-1$, the model is trained on all subsets except the i th subset, and evaluated on the i th subset. The average performance of the model across all k iterations is reported as the model accuracy.

The intended use of this method is to evaluate the difference in model performance as hyperparameters are varied, in order to choose the optimal hyperparameters for the context, without leaking information from the validation set. In this method, the final test data set is only consulted after the hyperparameters have been chosen. The implementation used in this assignment is shown below.

```

def kfold_cv(X, y, k, model):
    # Split the data into k folds
    X_folds = np.array_split(X, k)
    y_folds = np.array_split(y, k)
    accuracies = []
    for i in range(k):
        # Create training and validation sets
        # Training sets contain every fold except the ith
        X_train = np.concatenate(X_folds[:i] + X_folds[i+1:])
        y_train = np.concatenate(y_folds[:i] + y_folds[i+1:])
        # Validation set is the ith fold
        X_validation = X_folds[i]
        y_validation = y_folds[i]
        # Train the model
        model.fit(X_train, y_train)
        # Evaluate the model
        accuracies.append(model.score(X_validation, y_validation))
    return 1 - np.average(np.array(accuracies))

```

A small (and likely unrepresentative) experiment was conducted to choose a value for k to use in this assignment. A Logistic Regression model, using the default parameters was trained on the Fashion-MNIST dataset and scored against the test set; the score was computed using the k -Fold Cross Validation implementation above with k ranging from 5 to 10. The results of this experiment are shown in Figure 2. Since there was no significant difference in outcome between the values of k tested, the value of 5 was chosen for the remainder of the experiments to reduce runtime.

4 Logistic Regression

is a method for predicting discrete-valued labels based on input features which may or may not be discrete. In the case of this assignment the input vectors as 28x28 grayscale images with pixel values normalized to the range $[0, 1]$. The output labels are binary. Logistic Regression is used to learn (estimate) the underlying probability function which maps input vectors to output label likelihoods. The probability function is estimated using the maximum likelihood estimation.




Figure 2: LR Model Error vs. k

Logistic Regression can be thought of as a neural network in which the input features are multiplied by a weights matrix, the results are summed together and offset by a bias parameter.

$$z = \sum_{i=1}^n w_i x_i + b \quad (1)$$

Logistic Regression introduces a nonlinearity after the linear unit in the form of the sigmoid function to map the output of the linear model to the range $[0, 1]$. The sigmoid function is defined as

$$\sigma(z) = \frac{1}{1 + e^{-z}} \quad (2)$$

In binary classification, the output of the sigmoid function is then interpreted as the possibility that the input vector is an example of the positive class. The complementary probability represents the likelihood of the input example belonging to the negative case.

In the first experiment, the effects of the regularization parameter C on the Logistic Regression model were investigated. The regularization parameter is a hyperparameter which specifies the degree to which the model is

penalized for having large weights. In this experiment, higher values of C correspond to less regularization. The results of varying C are shown in Figure 3, note that the x-axis is presented on a logarithmic scale.




Figure 3: LR Model Error vs. C

The results of this experiment show that the model performs best when C is set to approximately 10^{-1} . High levels of regularization (the left side of the plot) show that high model error indicating that the model is underfitting the data as it is unable to learn the underlying probability function due to the restriction. It appears to be underfitting as both the training and testing error are quite high. On the right side of the plot, at much lower levels of regularization the model appears to show a classic pattern of overfitting where the training error continues to decrease while the testing error reaches a minimum and then begins to increase with decreasing regularization.

5 Support Vector Machines

Support Vector Machines are a supervised learning method which can be used in classification problems. The method attempts to find a hyperplane through the input space which separates the data and categorizes them correctly. The complexity of the hyperplane is controlled by the kernel used. In this experiment a linear kernel was used. The hyperplane is found by maximizing the margin between the hyperplane and the nearest data points on either side. During learning, a hyperplane which minimizes training error is learned.

This experiment varied the regularization parameter C in the same way as the previous experiment. The results of this experiment are shown in Figure 4. The results show a similar pattern as the previous, with high levels of regularization resulting in underfitting and low levels of regularization resulting in overfitting. In both cases the model error is quite high. The optimal value of C appears to be approximately 10^2 .




Figure 4: SVM Model Error vs. C

5.1 Finding the Optimal Regularization Parameter

For both the Logistic Regression and Support Vector Machine models, the regularization parameter C was varied in the range $[2^{-15}, 2^{11}]$. This range was chosen as it fully encompassed the range of values which produced the best results in the previous experiments. Both classifiers were trained on the training set, and scored using k-fold cross validation. The purpose of this section is to find the optimal value of C for each model without involving the test set. This represents a hyperparameter search before the final model weights are learned; so that the model can better approximate the true underlying distribution without overfitting to the test set.

The results of this experiment are shown in Figure 5. The bounds of the search are just wide enough to see the underfitting / overfitting pattern from the previous experiments. Both models reach an optimal value near the same point.

Logistic Regression Minimum Error:	0.15083
Optimal Logistic Regression C Value:	0.06250
SVM Minimum Error:	0.15291
Optimal SVM C Value:	0.03125

The two models were compared by computing the difference in their log loss scores on the test set. This calculation provides a more confident way of determining the better performing model. The confidence value is calculated according to:

$$\frac{1}{n} \sum_{i=1}^n [\ell(y_i, \hat{h}_1(x_i)) - \ell(y_i, \hat{h}_2(x_i))] \quad (3)$$

where

$$\ell(y, \hat{y}) = -y \log \hat{y} - (1 - y) \log(1 - \hat{y}) \quad (4)$$

The resulting confidence value was 2.212, a decidedly positive value indicating that the second hypothesis, in this case the SVM, is better than the first.




Figure 5: Model Error vs. Regularization, Using 5-fold Cross Validation

6 Gaussian Kernel SVM

The previous experiments used a linear kernel for the Support Vector Machine model. In this experiment a Gaussian kernel was used. Gaussian SVM models use a bandwidth parameter, γ to control the sensitivity of the kernel to the distance to near data points. Higher values of γ result in a more complicated decision surfaces which bend more radically to fit the data, and therefore is also more prone to overfitting. To combat overfitting for a given value of γ , the regularization parameter C was varied for each γ value tried. The result was to product a series of pairs, each containing a choice for γ and the correspondingly optimal choice for C .

The results of this experiment are shown in Figure 6. The horizontal axis represents the bandwidth parameter γ , and ther vertical represents the error of a model trained with the γ , C pair.

SVM Minimum Error:	0.14749
SVM Gamma Value:	0.00781
SVM C Value:	4.0




Figure 6: Model Error vs. Bandwidth Parameter, Using 5-fold Cross Validation

The results of this experiment show that the model performs best when γ is set to approximately 0.0078, with a correspondingly high regularization term of 4.0. Using a Gaussian kernel has improved the performance of the model; reducing the error from 15.29% to 14.75%.