

SENG 474 A02: Assignment 2

Model Selection Experiments: *Logistic Regression, Support
Vector Machines, and k-Fold Cross Validation*

Sean McAuliffe, V00913346

February 27th, 2023

Contents

1	Background	3
2	Environment and Implementation	3
3	k -Fold Cross Validation	3
4	Logistic Regression	5

1 Background

This report summarizes the results of a series of experiments conducted which implement two methods for binary classification: Logistic Regression, and Support Vector Machines (SVM). The purpose of this exercise was to compare the performance of these two methods, as well as to practice the use of cross-entropy loss and k -Fold Cross Validation for selection of optimal model hyperparameters.

For each of these approaches a number of models were trained on the Fashion-MNIST [dataset](#). The Fashion-MNIST dataset is a collection of 70,000 grayscale 28x28 pixel images of clothing items in 10 classes. For the purposes of these experiments, only classes 0 and 6 (T-shirt and Shirt) were used. Furthermore, to reduce runtimes during SVM training, only 2,400 training images, and 2000 testing images were used throughout the experiments.

2 Environment and Implementation

Linear Regression and SVM implementation from **scikit-learn** were used for this assignment. Of note for this report, the **scikit-learn** library implements the *linear_model.LinearRegression* and *svm.SVC* classes for Logistic Regression and SVM, respectively. These classes were used to train and test the models for this assignment.

All code for this assignment, including data import, preprocessing model creation, training, evaluation and plot generation was done in Jupyter Notebooks running on a **Python 3.10.6** kernel. Relying on utility functions provided with the Fashion-MNIST dataset, the source code for which has been provided with included with the .ipynb files with this report.

3 k -Fold Cross Validation

The k -Fold Cross Validation method is an improvement on the simple cross-validation which was used to evaluate model performance in the previous assignment. Cross validation involves splitting the available data into

training and validation sets; models will be trained on the training set, and scored against the validation set. In this way the true risk of the model can be estimated on data which was not used to train the model.

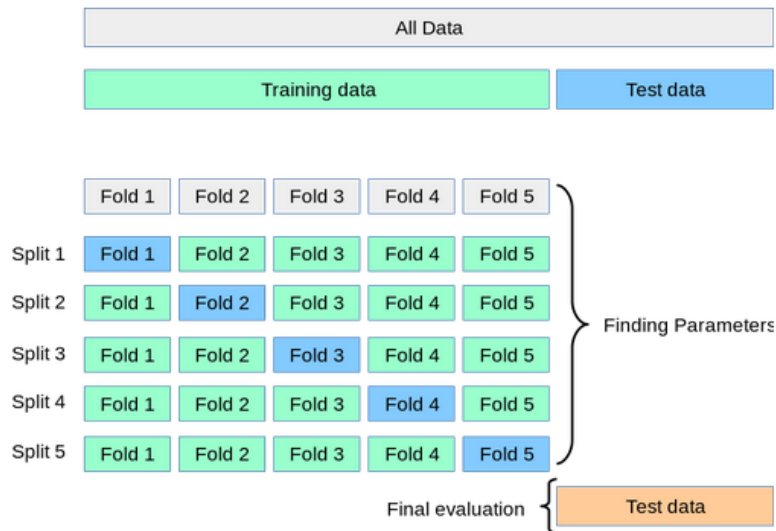


Figure 1: k-fold cross validation method, source scikit-learn.org

The k -Fold Cross Validation method is a better way to evaluate model performance. It is more robust to the choice of training and validation sets, and is more computationally efficient than the simple method. The k -Fold Cross Validation method involves splitting the available data into k equal sized subsets. Then for each i between 0 and $k-1$, the model is trained on all subsets except the i th subset, and evaluated on the i th subset. The average performance of the model across all k iterations is reported as the model accuracy.

The intended use of this method is to evaluate the difference in model performance as hyperparameters are varied, in order to choose the optimal hyperparameters for the context, without leaking information from the validation set. In this method, the final test data set is only consulted after the hyperparameters have been chosen. The implementation used in this assignment is shown below.

```

def kfold_cv(X, y, k, model):
    # Split the data into k folds
    X_folds = np.array_split(X, k)
    y_folds = np.array_split(y, k)
    accuracies = []
    for i in range(k):
        # Create training and validation sets
        # Training sets contain every fold except the ith
        X_train = np.concatenate(X_folds[:i] + X_folds[i+1:])
        y_train = np.concatenate(y_folds[:i] + y_folds[i+1:])
        # Validation set is the ith fold
        X_validation = X_folds[i]
        y_validation = y_folds[i]
        # Train the model
        model.fit(X_train, y_train)
        # Evaluate the model
        accuracies.append(model.score(X_validation, y_validation))
    return 1 - np.average(np.array(accuracies))

```

A small (and likely unrepresentative) experiment was conducted to choose a value for k to use in this assignment. A Logistic Regression model, using the default parameters was trained on the Fashion-MNIST dataset and scored against the test set; the score was computed using the k -Fold Cross Validation implementation above with k ranging from 5 to 10. The results of this experiment are shown in Figure 2. Since there was no significant difference in outcome between the values of k tested, the value of 5 was chosen for the remainder of the experiments to reduce runtime.

4 Logistic Regression




Figure 2: LR Model Error vs. k